

# Compiling in Parallel

Tools and Libraries for Compile-time Software Engineering  
HiPEAC Conference 2024  
Munich, Germany

Paul Keir<sup>1</sup>  
Joel FALCOU<sup>2</sup>

<sup>1</sup>School of Computing, Engineering & Physical Sciences  
University of the West of Scotland, Paisley, UK

<sup>2</sup>Le Laboratoire Interdisciplinaire des Sciences du Numérique (LISN)  
Université Paris-Saclay, Paris, France

January 17th, 2024

# Overview

- ▶ C++ Extensions for Parallelism
- ▶ Parallel Evaluation within the ClangOz Compiler
- ▶ A New Execution Policy Extension for the Standard Library
- ▶ Benchmark Program Analysis
- ▶ Related Work
- ▶ Conclusion and Future Work

# C++ Extensions for Parallelism

- ▶ Introduced to the standard *algorithms* library in C++17
- ▶ A set of overloads of existing standard library algorithms
- ▶ Each function template accepts a new *execution policy* parameter

```
constexpr std::execution::sequenced_policy      seq{};  
constexpr std::execution::parallel_policy       par{};  
constexpr std::execution::parallel_unsequenced_policy par_unseq{};  
constexpr std::execution::unsequenced_policy   unseq{}; // C++20
```

- ▶ The first `for_each` below increments all elements of `v` in serial
- ▶ The `std::execution::par` object can permit parallel execution
- ▶ Common serial execution is also available via `std::execution::seq`

```
std::vector<int> v{1, 2, 3, 4, 5, 6, 7, 8};  
std::for_each(v.begin(), v.end(), [](int& i) { i++; });  
std::for_each(std::execution::par, v.begin(), v.end(), [](int& i) { i++; });  
std::for_each(std::execution::seq, v.begin(), v.end(), [](int& i) { i++; });  
assert(v == (std::vector{4, 5, 6, 7, 8, 9, 10, 11}));
```

## Parallel *for loops*: the ClangOz Low-Level Ininsics API

The following conditions must be met for a *for loop* to be parallelised:

1. The targeted loop must be within a `constexpr` function;
2. The function must include an `execution::par` parameter;
3. The loop must come immediately after the ClangOz intrinsics:

# Parallel *for* loops: the ClangOz Low-Level Ininsics API

The following conditions must be met for a *for* loop to be parallelised:

1. The targeted loop must be within a `constexpr` function;
2. The function must include an `execution::par` parameter;
3. The loop must come immediately after the ClangOz intrinsics:

```
template <class T, class U>
constexpr
void __BeginEndIteratorPair(T& Begin, U& End);

template <class T, class U>
constexpr
void __PartitionUsingIndex(T LHS, U RHS, RelationalType RelTy);

template <class T>
constexpr
void __IteratorLoopStep(T& StartIter, OperatorType OpTy, const T& BoundIter);

template <class T>
constexpr
void __ReduceVariable(T Var, ReductionType RedTy, OperatorType OpTy);
```

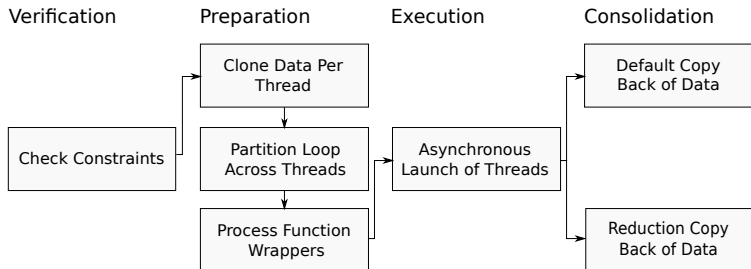
Compiler intrinsics: look like functions; no bodies; no-ops at runtime

# Limitations of the ClangOz compiler

- ▶ There is no support for nested parallelism
  - ▶ ...only the outer loop of a loop nest is parallelised;
- ▶ Only one loop can be parallelised within each function;
- ▶ Only containers owning contiguous data may be used;
  - ▶ **ContiguousContainer**  
std::string, std::vector, std::array, and built-in arrays (T[])
- ▶ Containers must be built upon pointer-based iterators.

# Parallelisation Phases within ClangOz

- ▶ `CallStackFrame` acts as a call stack for the constant evaluator
- ▶ `EvalInfo` holds information about the expression being evaluated

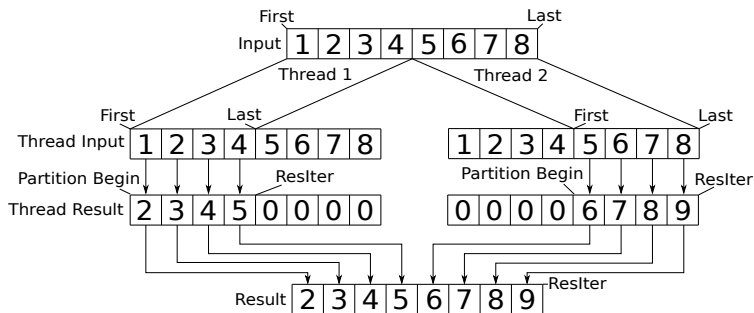


1. Check the 3 conditions are met; ensure no other active parallelism
2. `EvalInfo` clones; partition iteration space statically between threads
3. Launch, execute and await completion; last thread takes remainder
4. Synchronise thread data back into the main thread's `EvalInfo`

# Parallel Decomposition via PartitionedOrderedAssign

```
__BeginEndIteratorPair(First, Last);  
__IteratorLoopStep(ResIter, OperatorType::PreInc);  
__ReduceVariable(ResIter, ReductionType::PartitionedOrderedAssign,  
                 OperatorType::PreInc);
```

```
for (; First != Last; ++First) { *ResIter = (*First) + 1; ++ResIter; }
```



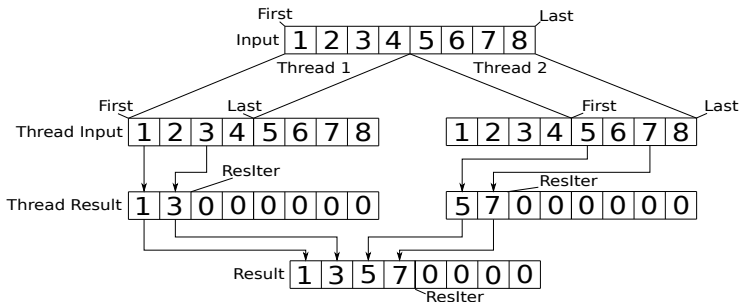
- ▶ `__IteratorLoopStep` asserts the loop/operator updates the pointer
- ▶ `PartitionedOrderedAssign` identifies a regular partitioning



# Parallel Decomposition via OrderedAssign

```
__BeginEndIteratorPair(First, Last);  
__ReduceVariable(ResIter, ReductionType::OrderedAssign,  
                 OperatorType::PreInc);
```

```
for (; First != Last; ++First)  
    if (*First % 2) { *ResIter = *First; ++ResIter; }
```



- ▶ Here an irregular partitioning is employed via `OrderedAssign`
- ▶ Useful when each thread creates distinct quantities for aggregation

# Implementing `for_each` using the Low-Level API

- ▶ But we don't want end users working with the low-level API
- ▶ Instead, we react to the *execution policy*: `std::execution::par`
  - ▶ ...when evaluated within a constant expression

# Implementing `for_each` using the Low-Level API

- ▶ But we don't want end users working with the low-level API
- ▶ Instead, we react to the *execution policy*: `std::execution::par`
  - ▶ ...when evaluated within a constant expression

Below we implement the standard function template `std::for_each`:

```
template <class _ExecutionPolicy, class _ForwardIterator, class _Function>
constexpr
__enable_if_constexpr_par_execution_policy_t<_ExecutionPolicy, void>
for_each(_ExecutionPolicy&& __exec, _ForwardIterator __first,
        _ForwardIterator __last, _Function __f)
{
    __BeginEndIteratorPair(__first, __last);
    __ReduceVariable(__first, PartitionedOrderedAssign, PreInc);

    for (; __first != __last; ++__first)
        __f(*__first);
}
```

We have implemented 30 function templates from the C++ standard Algorithms and Numerics libraries

## Example Usage of the High-Level API

```
constexpr bool g() {  
    std::vector<int> v{1, 2, 3, 4, 5, 6, 7, 8};  
    std::for_each(std::execution::par, v.begin(), v.end(),  
                  [](int &i) { i++; });  
    assert(v == (std::vector{2, 3, 4, 5, 6, 7, 8, 9}));  
}  
static_assert(g() == true);
```

- ▶ The function `g` increments each element of `v` in parallel
- ▶ The `static_assert` invokes constant expression evaluation
- ▶ Other execution policy overloads are not `constexpr`
- ▶ Traditional `std::for_each` is `constexpr`; but serial only

## Example Usage of the High-Level API

```
constexpr bool g() {  
    std::vector<int> v{1, 2, 3, 4, 5, 6, 7, 8};  
    std::for_each(std::execution::par, v.begin(), v.end(),  
                  [](int &i) { i++; });  
    assert(v == (std::vector{2, 3, 4, 5, 6, 7, 8, 9}));  
}  
static_assert(g() == true);
```

- ▶ The function `g` increments each element of `v` in parallel
- ▶ The `static_assert` invokes constant expression evaluation
- ▶ Other execution policy overloads are not `constexpr`
- ▶ Traditional `std::for_each` is `constexpr`; but serial only
  - ▶ ...though C'est 2 does include some `constexpr` `std::execution::seq` overloads

# Benchmarking Information

- ▶ Intel Core i9-12900K CPU; 32 GB RAM; 1 TB SSD
- ▶ 8 performance cores; 16 hyper-threads
- ▶ 64-bit Ubuntu under WSL2 virtual machine on Windows 11
- ▶ Executed using 2, 4, 8 and 16 threads; times averaged over 6 runs
- ▶ Times correspond to one or more consecutive `for_each` invocations


The Benchmark Programs:

- ▶ **Mandelbrot**\* Complex value iteration; max 128 iterations
- ▶ **Blackscholes**\*\* processes financial data using a PDE
- ▶ **N-Body**\* Jovian planets simulation via a symplectic-integrator

\* From The Computer Language Benchmarks Game

\*\* From Princeton's PARSEC Benchmark Suite

# Overall Speedup



# Related Work

*Optimizing gpu programs by partial evaluation*

A. Tyurin, D. Berezun, and S. Grigorev

Proceedings of 25th ACM PPoPP (2020)

*Anydsl: A partial evaluation framework for programming high-performance libraries*

R. Leißa, et al.

Proceedings of ACM on Programming Languages, OOPSLA (2018)

*Practical partial evaluation for high-performance dynamic language runtimes*

T. Würthinger et al.

Proceedings of the 38th ACM SIGPLAN Conference on PLDI (2017)

*Large-scale parallelization of partial evaluations in evolutionary algorithms for real-world problems*

A. Bouter, T. Alderliesten, A. Bel, C. Witteveen, and P. A. Bosman

Proceedings of the Genetic and Evolutionary Computation Conference (2018)

*Distributed partial evaluation*

M. Sperber, P. Thiemann, and H. Klaeren

Proceedings of the 2nd International Symposium on Parallel Symbolic Computation (1997)

*Partial evaluation in parallel*

C. Consel and O. Danvy

Lisp and Symbolic Computation (1993)



# Conclusion

- ▶ ClangOz, the first compiler capable of parallel evaluation of C++ constant expressions
- ▶ Support (subset) for the C++ Standard Algorithms & Numerics library
  - ▶ ...via the `std::execution::par` policy overloads
- ▶ A flexible, low-level intrinsics API
- ▶ A suite of `constexpr` benchmark programs, with study results

<https://github.com/agozillon/ClangOz>

## Future Work

- ▶ Add support for more C++ standard library functions
- ▶ Profile, to improve parallel performance

# Hands On

- ▶ wordcount; todo

# Acknowledgements

**UWS** UNIVERSITY OF THE  
WEST *of* SCOTLAND

**RSE** *The Royal Society  
of Edinburgh*  
KNOWLEDGE MADE USEFUL