

CTBENCH - COMPILE-TIME BENCHMARKING

Jules Pénuchot

January 15, 2024

PhD Student Parallel Systems, LISN, Paris-Saclay University



Metaprogramming is evolving

Support libraries

- Boost.Mpl
- Boost.Fusion
- Boost.Hana
- Boost.Mp11
- Brigand
- *that code snippet repository you probably own*

Applications

- Eigen - <https://eigen.tuxfamily.org/>
- Blaze - <https://bitbucket.org/blaze-lib/blaze>
- CTRE - <https://github.com/hanickadot/compile-time-regular-expressions>
- CTPG - <https://github.com/peter-winter/ctpg>

Metaprogramming lacks tooling

Metaprogramming is *almost* on par with regular programming...

- ...but regular programming has debuggers, profilers,
- We know how to benchmark it to get **meaningful**, quantitative results,
- No such process for meta-programs,
- Little to no **science** behind compile time rule of thumbs.
- **We need a sane process for understanding compile times.**

How to measure compile times?

- Templight, *Zoltán Borók-Nagy, Zoltán Porkoláb, and József Mihalicza* (2009)
- Metabench, *Louis Dionne and Bruno Dutra* (2016)
- Build-Bench, *Fred Tingaud* (2017)
- Clang time-trace & Clang Build Analyzer, *Aras Pranckevičius* (2019)

Introducing ctbench

What is ctbench?

<https://github.com/jpenuchot/ctbench>

- Compile time benchmarking & **data analysis** tool for Clang,
- built on-top of **time-trace**,
- **repeatability** & accuracy in mind,
- **variable size** benchmarks,
- C++ developer friendly:
 - C++ only benchmark files,
 - CMake API,
 - JSON config files (with a few ones already provided)

But how does it work?

Benchmarking methodology

- Benchmark set:
 - collection of benchmark cases to compare
 - Benchmark case:
 - compilable C++ file,
 - compiled several times for a given **range of sizes**,
 - benchmark **iteration size** passed as a preprocessor define
 - Benchmark iteration:
 - terminology for a benchmark case compiled with a **given size**,
 - **several samples** for each iteration size for improved accuracy
 - Sample:
 - **one** time-trace file
- Benchmark set → Benchmark cases → Benchmark iterations → Samples

Benchmark case

Recursive sum

- *ctbench* defines `BENCHMARK_SIZE` for each iteration size.

```
1  // Metaprogram to benchmark:
2  template <unsigned N> struct ct_uint_t { static constexpr unsigned value = N; };
3
4  template <typename T> auto sum(T const &) { return T::value; }
5  template <typename T, typename... Ts> auto sum(T const &, Ts const &...tl) {
6      return T::value + sum(tl...);
7  }
8
9  // Benchmark driver:
10 #include <boost/preprocessor/repetition/enum.hpp>
11 #define GEN_MACRO(Z, N, TEXT) TEXT<N> {}
12 unsigned foo() {
13     // return sum(ct_uint_t<1>{}, ..., ct_uint_t<BENCHMARK_SIZE>{});
14     return sum(BOOST_PP_ENUM(BENCHMARK_SIZE, GEN_MACRO, ct_uint_t));
15 }
```

Benchmark case

Expansion sum

- *ctbench* defines `BENCHMARK_SIZE` for each iteration size.

```
1  // Metaprogram to benchmark:
2  template <unsigned N> struct ct_uint_t { static constexpr unsigned value = N; };
3
4  template <typename... Ts> auto sum(Ts const &...) { return (Ts::value + ...); }
5
6  // Benchmark driver:
7  #include <boost/preprocessor/repetition/enum.hpp>
8  #define GEN_MACRO(Z, N, TEXT) TEXT<N> {}
9  unsigned foo() {
10     // return sum(ct_uint_t<1>{}, ..., ct_uint_t<BENCHMARK_SIZE>{});
11     return sum(BOOST_PP_ENUM(BENCHMARK_SIZE, GEN_MACRO, ct_uint_t));
12 }
```


- Benchmark declaration

```
ctbench_add_benchmark(variadic_sum.recursive # Target name
  variadic_sum/recursive.cpp                # Benchmark file
  1 32 1                                    # Start, stop, and step
  10)                                        # Number of repetitions
```

- Graph declaration

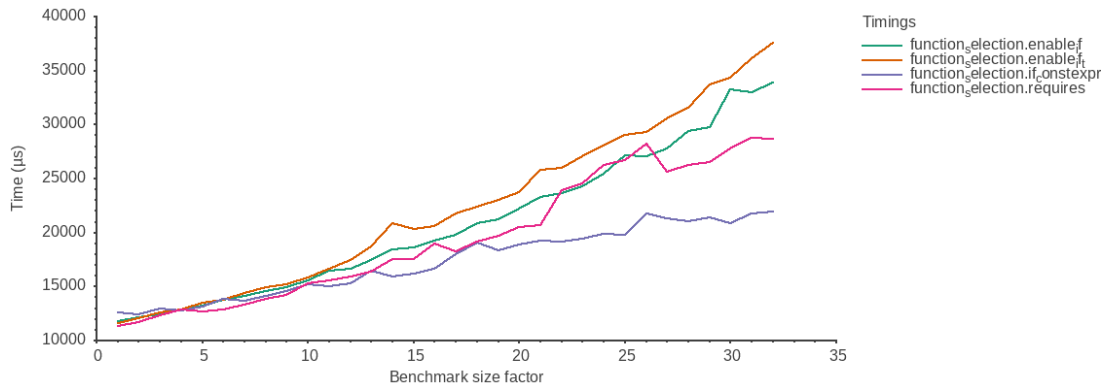
```
ctbench_add_graph(variadic_sum-graph # Target name
  configs/feature_comparison.json    # Config file
  variadic_sum.expansion              # Benchmark target
  variadic_sum.recursive)            # ...
```

- Optional: Bring your own flags with `ctbench_add_custom_benchmark`

Sample benchmarks

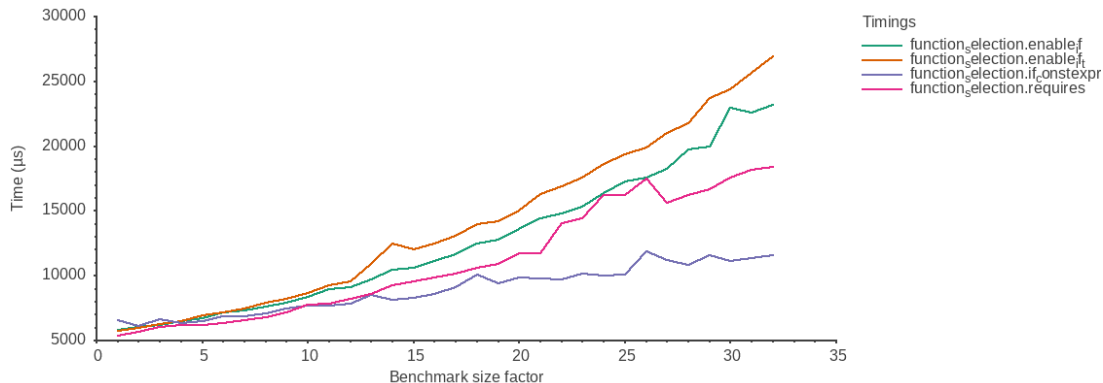
Function selection

- `enable_if_t`, `enable_if`
- `if constexpr`
- `requires`



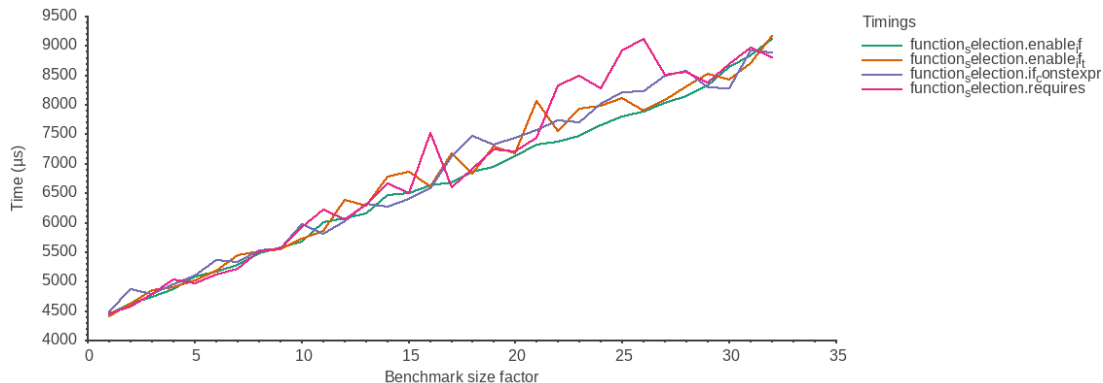
Function selection

- `enable_if_t`, `enable_if`
- `if constexpr`
- `requires`



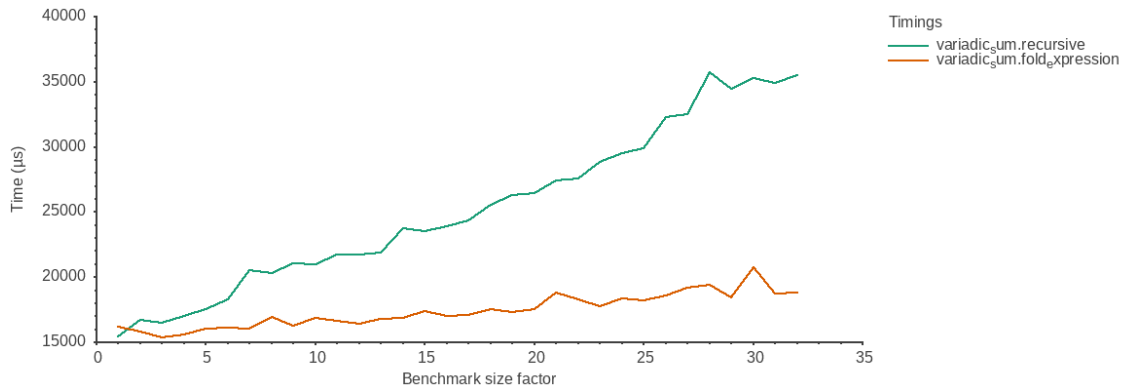
Function selection

- `enable_if_t`, `enable_if`
- `if constexpr`
- `requires`



Variadic sum

- recursive
- fold_expression




Targeted data: ExecuteCompiler

C++ contenders

Variadic sum

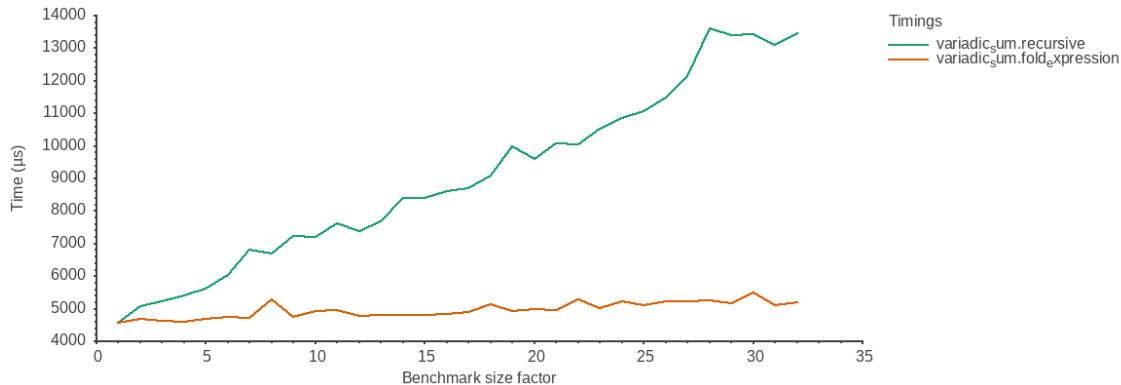
- recursive
- fold_expression



Targeted data: Frontend

Variadic sum

- recursive
- fold_expression



Targeted data: Backend

Conclusion

Project overview

- CMake API
 - **benchmarking.cmake** declares the end-user API,
 - Documentation is provided inside (easily extracted into a MD file)
- **grapher** subproject (meatiest part)
 - CLI, time-trace file reading, predicate engine, and plotting,
 - Designed as a **library** + CLI drivers,
 - relies heavily on **Sciplot** (<https://github.com/Sciplot/Sciplot>),
 - new plotters can be written easily
- Tooling:
 - **time-trace-wrapper**: clang exec wrapper to extract time-trace files
 - **cmake-doc-extractor**: extracts the API doc into a MD file

Thanks for your attention !